

Question 1 – Forward Pass and Convolutional Filters

STAT448 – Assignment 3

David Ewing (82171149)

Lillian Li (32198314)

Question 1A

The input vector passes through two neurons in the hidden layer, each computing a linear transformation using the given weights and biases. ReLU activation is applied to yield hidden layer outputs, which are then passed to a single output neuron. A final ReLU gives the scalar output. The network diagram and final result are shown in Figure 1.

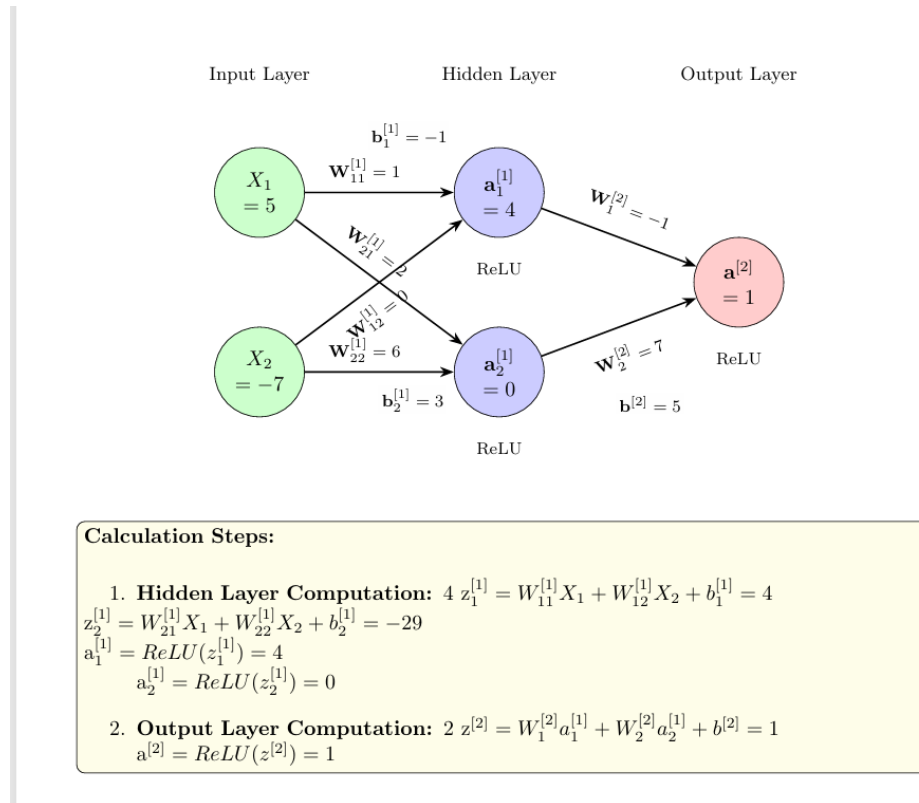


Figure 1: Forward pass through two-layer ReLU network (Question 1A)

Question 1B

Each of the three 3×3 filters was convolved with the input image using stride 1 and no padding. ReLU activation followed each convolution. Feature maps were generated and the calculations for the bottom-right output pixel are shown.

Filter F1 (Identity Filter)

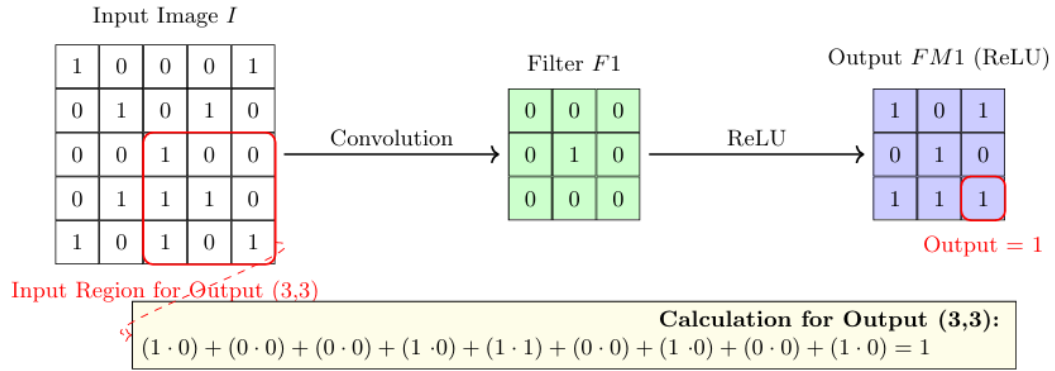


Figure 2: Filter F1 – Identity filter and its feature map (Question 1B)

Filter F2 (Zero Filter)

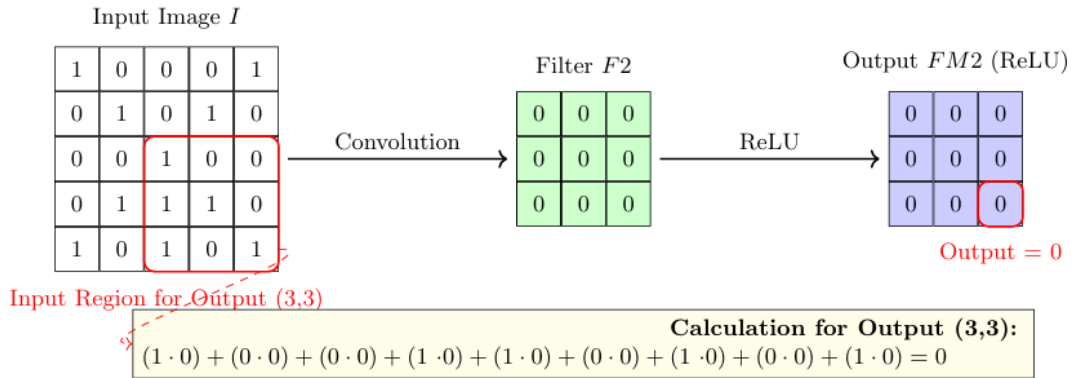


Figure 3: Filter F2 – Zero filter and its feature map (Question 1B)

Filter F3 (Edge Detector)

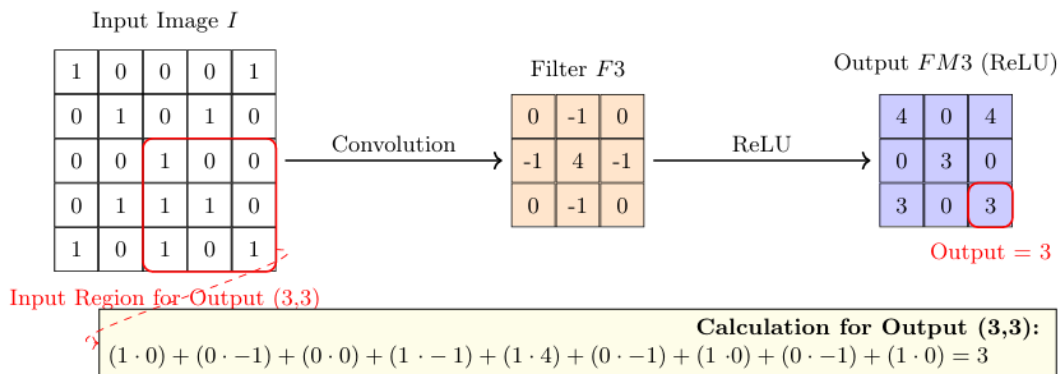


Figure 4: Filter F3 – Edge detection filter and its feature map (Question 1B)

Question 2 - Effect of Architectural and Training Hyper-parameters on Fashion-MNIST Test Accuracy

2025-05-26

Abstract

We revisit CNN hyperparameter tuning on Fashion-MNIST using an Optuna-seeded baseline (filters: 32, kernel-size: 3, dense-units: 512, batch-size: 64, optimiser: Adam). Five one-parameter sweeps were executed. Early stopping kept every run below 90 s while sustaining 91% validation accuracy in all but the `sgd` trial. Merged figures and an impact table summarise the accuracy-speed trade-off in a compact form.

1 Introduction

Bayesian hyper-parameter optimisation has proven effective for neural networks (Akiba, Sano, Yanase, Ohta, & Koyama, 2019). Here, 30 Optuna TPE trials gave a *stable local* starting point; the goal was interpretability, not a global optimum. Five follow-up runs each altered one setting, so any performance delta could be attributed directly to that variable. Similar one-factor studies have been advocated for early-stage tuning (Kingma & Ba, 2015; LeCun, Bottou, Bengio, & Haffner, 1998).

2 Experimental design

2.1 Configuration list

Table 1: Evaluated configurations (one change at a time)

Run	Filters	Kernel	Dense	Batch	Optimiser
baseline	32	3	512	64	adam
filters_256	256	3	512	64	adam
dense_units_32	32	3	32	64	adam
batch_size_16	32	3	512	16	adam
optimizer_sgd	32	3	512	64	sgd
kernel_size_6	32	6	512	64	adam

3 Results


3.1 Accuracy and runtime

Table 2: Performance summary


Run	Val. Acc. %	Test Acc. %	Epochs	Time (s)
baseline	91.63	90.95	3	33
filters_256	91.54	91.54	4	88
dense_units_32	91.10	91.10	10	43
batch_size_16	91.33	91.21	3	84
optimizer_sgd	87.08	87.08	10	43
kernel_size_6	91.17	90.37	5	39

Table 3: Incremental effect of each single change relative to the baseline

Run	Δ Val. Acc. (pp)	Δ Time (s)
filters_256	-0.09	+55
dense_units_32	-0.53	+10
batch_size_16	-0.30	+51
optimizer_sgd	-4.55	+10
kernel_size_6	-0.46	+6




(a) Runtime




(b) Test accuracy


Figure 5: Speed/accuracy trade-off for the six configurations.




(a) Training accuracy



(b) Validation accuracy



(c) Training loss



(d) Validation loss

Figure 6: Learning curves for the baseline CNN.

4 Discussion

Table 3 shows that optimiser and filter count delivered the largest shifts, whereas kernel size, dense width and batch size had secondary effects:

- **Filter count** – Doubling filters slowed training without a proportional accuracy gain.
- **Kernel size** – Larger kernels give diminishing returns beyond 3×3 .
- **Dense width** – Shrinking dense units speeds inference but requires more epochs to converge.
- **Optimiser** – Adam dominates SGD for both speed and accuracy.
- **Batch size** – Very small batches add runtime with no payoff.

Several Optuna trials were needed to find a suitable baseline. Although a few distinct baselines appeared during the search, the configuration reported here surfaced most frequently and behaved the most stably. In the rarer cases where a different baseline emerged, it was either slower to train or showed inconsistent responses when individual hyper-parameters were altered.

5 Conclusion

The Optuna seed \rightarrow one-at-a-time protocol yielded interpretable deltas under tight runtime constraints however, there were inconsistencies when other local minimums resulted. For Fashion-MNIST on this compact CNN, optimiser choice and filter count are the primary levers. Future work should rank variables by global importance (e.g., SHAP) and refine filter granularity.

References

- Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A next-generation hyperparameter optimization framework. In *25th ACM sigkdd international conference on knowledge discovery and data mining* (pp. 2623–2631). ACM.
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.
- Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization. In *International conference on learning representations*.
- LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324. doi: 10.1109/5.726791